

```
[1]: %matplotlib inline
import matplotlib

# make figures bigger:
import numpy as np
import matplotlib.pyplot as plt
import random
import scipy, scipy.stats

# make figures better:
font = {'weight': 'normal', 'size': 20}
matplotlib.rc('font', **font)
matplotlib.rc('figure', figsize=(8.0, 6.0))
matplotlib.rc('xtick', labelsiz=16)
matplotlib.rc('ytick', labelsiz=16)
matplotlib.rc('legend', **{'fontsize': 16})
matplotlib.rc('image', cmap='viridis') # change default colormap
```

## DS1 Lecture 24

James Bagrow, james.bagrow@uvm.edu, <http://bagrow.com>

---

### Last time:

1. Intro to Bayesian Inference
  - Checking whether a coin is fair
2. Using posterior distributions for model selection: Bayes Factor
  - Unbiased (simple) model vs. Biased (complex) model (online)

### Today:


1. Bayesian inference
  - text messaging data

## Bayesian inference

We have some data from a cell phone user, the number of text messages received per day, during a time period:

```
[2]: C = np.loadtxt("txtdata.csv")
T = len(C)

plt.bar(np.arange(T), C, color='C0', ec='none')
plt.xlabel("Time (days)")
plt.ylabel("Text messages received")
plt.xlim(0, T)
plt.show()
```



We want to see: Does the **rate** (messages per day) over the course of the data?

- If it does change, when does it change? Does it change suddenly or gradually?

Visually, it looks pretty steady...

Here's a quick calculation, mean rate of messages averaged over the beginning and ending of the data:

```
[3]: print("mean rate during first month =", C[:30].mean(), "msgs/day")
      print("mean rate during last month  =", C[-30:].mean(), "msgs/day")
```

mean rate during first month = 19.9 msgs/day

mean rate during last month = 22.7 msgs/day


And let's visualize these rates:

```
[4]: plt.bar(np.arange(T), C, color='C0', ec='none')

      # add mean rate for first and last 30 days:
      n = 30
      plt.plot(np.arange(n),          n*[C[:n].mean() ], 'r-', lw=4, alpha=0.667 )
      plt.plot(np.arange(T-n+1,T+1), n*[C[-n:].mean()], 'r-', lw=4, alpha=0.667 )

      plt.xlabel("Time (days)")
      plt.ylabel("Text messages received")
```

```
plt.xlim(0-0.5, T)
plt.show()
```



Let's zoom in a bit:


```
[5]: plt.bar(np.arange(T), C, color='C0', ec='none')

# add mean rate for first and last 30 days:
n = 30
plt.plot(np.arange(n), n*[C[:n].mean() ], 'r-', lw=4, alpha=0.667 )
plt.plot(np.arange(T-n+1,T+1), n*[C[-n:].mean()], 'r-', lw=4, alpha=0.667 )

# and also, add a linear model!
slope, intercept, r_value, p_value, std_err = scipy.stats.linregress(np.arange(T), C)
plt.plot(np.arange(T), slope*np.arange(T)+intercept, 'g--')
print("p-value (linear model) =", p_value)

plt.xlabel("Time (days)")
plt.ylabel("Text messages received")
plt.xlim(0-0.5, T)
plt.ylim(0,30) # ***
plt.show()
```

p-value (linear model) = 0.7801884386142313



Hmm,  $\approx 3$  messages per day difference... Hmm... ....

## Statistical model for our hypothesis

[notes on board about statistical model, poisson distributions, rates, proof of poisson as limit of binomial]

Let's **simulate** this thing.

First, a function that generates data from the model:

```
[6]: def gen_artificial_sms_dataset(lambda_1, lambda_2, tau, length=80):  
    """vector of counts drawn from Pois(lambda_1) when t < tau  
    and Pois(lambda_2) when t >= tau.  
    """  
    data_b = np.random.poisson(lam=lambda_1, size=tau)  
    data_a = np.random.poisson(lam=lambda_2, size=length-tau)  
  
    return np.append(data_b, data_a)
```

Now let's generate **artificial text data** following our model. We'll pick some values of  $\lambda_1$ ,  $\lambda_2$ , and  $\tau$ , and make a plot for each:

```

[7]: # specify FOUR models (lam1, lam2, tau):
models = [
    (50.01, 33.02, 51),
    (23.87, 7.02, 19),
    ( 6.19, 16.40, 26),
    ( 8.70, 3.99, 12)
]

# set up a figure with 4 plots stacked in a row:
fig, list_axes = plt.subplots(4, 1, sharex=True)
fig.set_size_inches(8, 10)
fig.subplots_adjust(hspace=0.1)


for ax, (L1, L2, tau) in zip(list_axes, models):
    data = gen_artificial_sms_dataset(L1, L2, tau)

    ax.bar(np.arange(len(data)), data, color='C0', ec='none')
    ax.bar(tau-1, data[tau-1], color='C1', ec='none')

    txt = ax.text(0.725, 0.85,
        "$\lambda_2/\lambda_1 = {:.3f}$".format(L2/L1),
        fontsize=18,
        horizontalalignment='left',
        verticalalignment='center',
        transform=ax.transAxes
    )
    txt.set_bbox(dict(color='white', alpha=0.5))

# clean up the plots:
plt.xlabel("Day")
plt.xlim(-0.5, 80)
fig.text(0.06, 0.5, 'Number of texts',
    fontsize=22,
    ha='center', va='center',
    rotation='vertical'
);

```



Of course, we aren't GIVEN the parameters  $\lambda_1$ ,  $\lambda_2$ , and  $\tau$ . We need to infer them. How? Bayes!

### Posterior!

$$\Pr(\text{model} \mid \text{data}) = \Pr(\text{data} \mid \text{model}) \Pr(\text{model}) / \Pr(\text{data})$$

We want to get  $\Pr(M \mid D)$  to find which parameters ( $\lambda_1$ ,  $\lambda_2$ , and  $\tau$ ) have high probability given the observed data.

If we are **given values** of the parameters we can compute:

Likelihood \* Prior = Pr(data | model) Pr(model) ~ Posterior

This is **almost** what we need—we are missing  $P(\text{data})$ . Next time we will fill in the gap to perform inference.

---

For now let's make some pictures to see priors and posteriors. \* Simplify: We will fix  $\tau$  and only look at  $\lambda_1$  and  $\lambda_2$ .

## Priors

We will look at two priors for the  $\lambda$ 's

### Uniform Prior:

All models (values of  $\lambda$ 's) over a certain range are equally likely before we look at the data.

- $\lambda_1 \sim \text{Uniform}(0,5)$
- $\lambda_2 \sim \text{Uniform}(0,5)$

(This is simpler than the prior we used on the board.)


Plot the prior:

```
[8]: # double-loop of 100 values for lambdas:
lam1 = lam2 = np.linspace(0, 5, 100) # vectors
L1, L2 = np.meshgrid(lam1, lam2)      # for 3d plots

# compute probabilities:
Uniform_Prior = 0.2*0.2*np.ones(L1.shape)
#uni_L1 = scipy.stats.uniform.pdf(L1, loc=0, scale=5)
#uni_L2 = scipy.stats.uniform.pdf(L2, loc=0, scale=5)
#Prior = np.dot(uni_L1[:, None], uni_L2[None, :])

# top-down (matrix) view:
plt.imshow(Uniform_Prior, interpolation='none', origin='lower',
           vmax=1, vmin=-0.15,
           extent=(0, 5, 0, 5)
           );

plt.xlim(0, 5); plt.xlabel("$\lambda_1$")
plt.ylim(0, 5); plt.ylabel("$\lambda_2$")
plt.title(r"Uniform prior, no $\tau$.");
```



(In practice we wouldn't make such a visualization, and we would probably want to include a *colorbar*, but here's it's meant to compare with what comes next.)

Boring... Here is a 3D view:

```
[9]: from mpl_toolkits.mplot3d import Axes3D


# 3D view:b
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')

ax.plot_surface(L1, L2, Uniform_Prior, cmap=plt.cm.viridis)
ax.set_xlabel("\n$\lambda_1$")
ax.set_ylabel("\n$\lambda_2$");

ax.zaxis.set_rotate_label(False) # disable automatic rotation
ax.set_zlabel("\nPrior prob.", linespacing=-3, rotation=-90)
```



```
ax.dist = 11 # camera distance
```



Also pretty boring...

### Exponential Prior:

Here we now suppose that smaller values of  $\lambda$  are “more likely” than larger values, all else being equal.

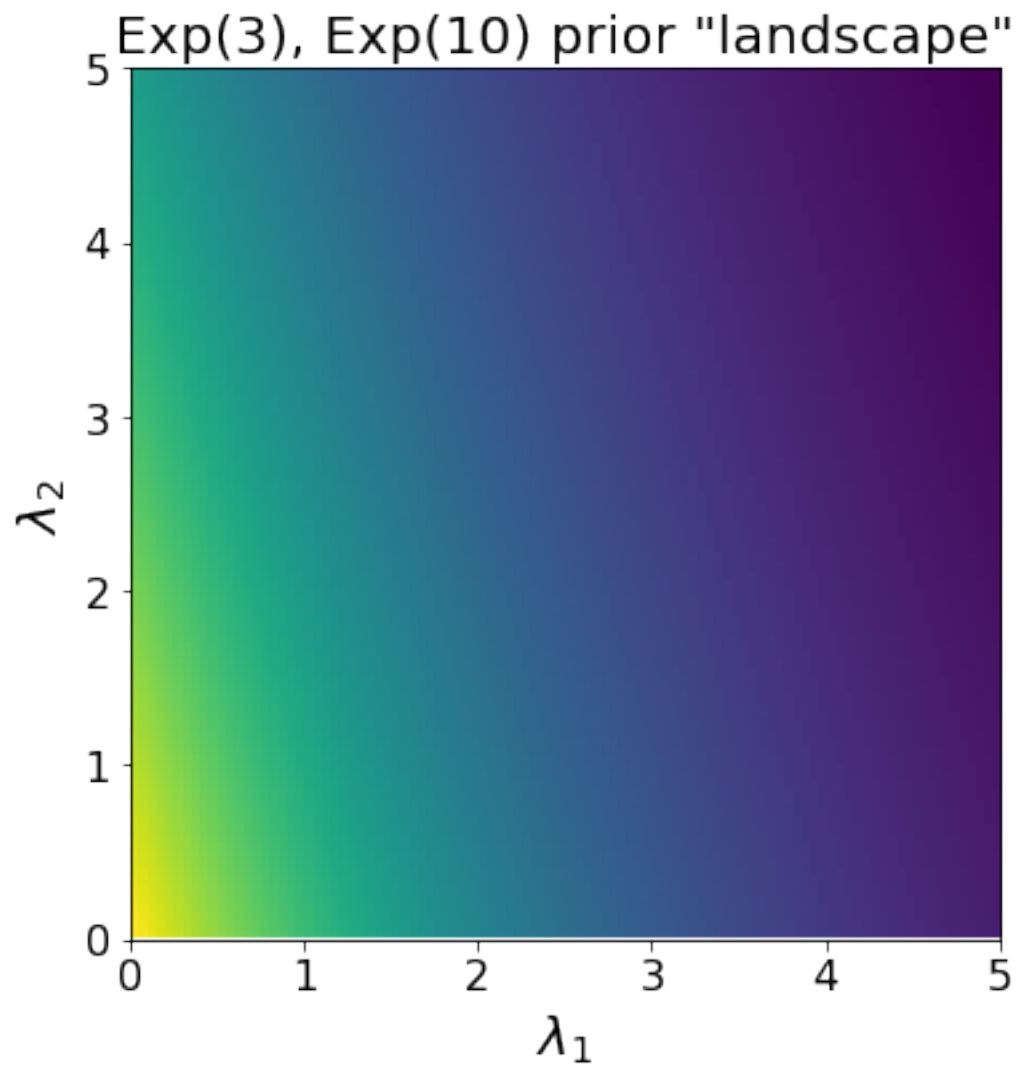
Example:

- $\lambda_1 \sim \text{Exp}(1/3)$      $\text{mean}(\lambda_1) = 3$
- $\lambda_2 \sim \text{Exp}(1/10)$      $\text{mean}(\lambda_2) = 10$

```
[10]: exp_L1 = scipy.stats.expon.pdf(lam1, scale=3)
exp_L2 = scipy.stats.expon.pdf(lam2, scale=10)
Exp_Prior = np.dot(exp_L2[:, None], exp_L1[None, :])

# plot:
im = plt.imshow(Exp_Prior, interpolation='none', origin='lower',
                extent=(0, 5, 0, 5))

plt.title("Exp(3), Exp(10) prior \"landscape\"", fontsize=20)
plt.xlabel("$\lambda_1$")
plt.ylabel("$\lambda_2$");
```




(Like before, needs a colorbar.)

And here the 3d view is not so boring:

```
[11]: fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.plot_surface(L1, L2, Exp_Prior, cmap=plt.cm.viridis)
ax.view_init(azim=390)

ax.set_xlabel("\n$\lambda_1$")
ax.set_ylabel("\n$\lambda_2$")

ax.set_zlabel("Prior prob.", labelpad=11);
```



Where does the knowledge of the prior come from? Past experience, physical reasons, etc.

- Uniform prior: we have no way “a priori” to assume any value of the parameter over another value.
- Exponential prior: we think both parameters are more likely to come from small values than from large values.

## Posteriors

**Now comes the data!**

The data *warp*s the height of the probability surface via the likelihood. This is Bayes Theorem.

Let’s make some fake data from a (secretly known) model. We will at first draw a few values and see how the posterior changes for both priors.

```
[12]: # create the observed data

# sample size of data we observe, trying varying this (keep it less than 100 ;)
N = 3 # 2N total datapoints, N on either side of tau
# the true parameters, but of course we do not see these values...
lambda_1_true = 1
lambda_2_true = 3

# Generate count data, given the above true lambdas:
# (could also have used gen_artificial_sms_dataset...)
```

```
data_b = scipy.stats.poisson.rvs(lambda_1_true, size=(N, 1))
data_a = scipy.stats.poisson.rvs(lambda_2_true, size=(N, 1))
data = np.append(data_b, data_a) # this is our synthetic c_t
```

And generate the likelihoods for each pair of  $\lambda$ 's:

```
[13]: # parameter sweep over lambdas to get likelihoods...
L1 = L2 = np.linspace(0.01, 5, 100)

Pois = scipy.stats.poisson.pmf # Pois(x,L) = L^x exp(-L) / x!

like_db = np.array([Pois(data_b, lam1).prod() for lam1 in L1])
like_da = np.array([Pois(data_a, lam2).prod() for lam2 in L2])
likelihood = np.dot(like_da[:, None], like_db[None, :]) # matrix
```

```
[14]: if N < 5:
    print(data)
    print(Pois(data, 0.3))
    print()
    print(Pois(data, 0.3).prod())
```

```
[0 2 1 2 1 4]
[7.40818221e-01 3.33368199e-02 2.2245466e-01 3.33368199e-02
 2.2245466e-01 2.50026149e-04]
```

```
1.0167431302704622e-08
```

Put it into a little *reusable* function:

```
[15]: def gen_data_likelihood(N):
    data_b = scipy.stats.poisson.rvs(lambda_1_true, size=(N, 1))
    data_a = scipy.stats.poisson.rvs(lambda_2_true, size=(N, 1))
    data = np.append(data_b, data_a)

    like_db = np.array([Pois(data_b, lam1).prod() for lam1 in L1])
    like_da = np.array([Pois(data_a, lam2).prod() for lam2 in L2])
    likelihood = np.dot(like_da[:, None], like_db[None, :]) # matrix

    return data, likelihood
```

Now, plot:

```
[16]: def plot_mat(X,Y,Z, fig=None, extent=None):
    if fig is None:
        fig = plt.gcf() # get current fig
    ax = fig.gca()
    ax.imshow(Z, interpolation='none', origin='lower',
              extent=extent)
    if extent is not None:
        ax.set_xlim(extent[0:2])
        ax.set_ylim(extent[2:] )

def plot_truth(x_true,y_true, fig=None):
```

```

if fig is None:
    fig = plt.gcf() # get current fig
ax = fig.gca()
ax.scatter(x_true, y_true, c="w", s=50, edgecolor="k",zorder=1e9)

```

```

[17]: def each_plot(subplot, Y, title=None):
    plt.subplot(subplot)
    plot_mat(L1, L2, Y, extent=extent)
    plot_truth(lambda_1_true, lambda_2_true)
    plt.title(title, fontsize=18)

def label_axes():
    fig.text(0.5, 0.06, r"$\lambda_1$",
            fontsize=22,
            ha='center', va='center')
    fig.text(0.06, 0.5, r"$\lambda_2$",
            fontsize=22,
            ha='center', va='center',
            rotation='vertical');

```

Now, *really* plot:

```

[21]: fig = plt.figure(figsize=(9,9))
    extent = (0,5,0,5) # = (xmin,xmax,ymin,ymax)


    N = 1000 # 10,20,100,200,1000 -- there are 2N datapoints
    data, likelihood = gen_data_likelihood(N)

    each_plot(221, Uniform_Prior, title='Uniform prior')
    each_plot(223, Uniform_Prior * likelihood, title=r"$\propto$ Posterior (Uni.)")

    each_plot(222, Exp_Prior, title=r"Exponential Prior")
    each_plot(224, Exp_Prior * likelihood, title=r"$\propto$ Posterior (Exp.)")

    label_axes()

```



Now let's go back and see how things look when there's more data! \* [Rerun for different values of  $N$ ]

Remarks:

- These functions are only *proportional* to the posterior because we have ignored the denominator in Bayes Thm.
- The posterior becomes more sharply peaked as  $N$  increases
- Sometimes the “bulk” of the posterior will not include the secretly known true value  $(\lambda_1, \lambda_2)$ , but it will be close to  $(\lambda_1, \lambda_2)$ .
- The uniform and exponential priors look very different, but their corresponding posteriors look very similar
- The real *cheat* here is the absence of  $\tau$ .

Bayesian inference. Take your (meaningful?) priors, update them with the likelihoods (using the observed data) to get a function proportional to the posterior.

- Of course, **parameter sweeps** like we've done (looping over many combinations of  $\lambda_1$  and  $\lambda_2$ ) are prohibitive in practice. Instead, **drawing samples** from the posterior gives us distributions of our parameters that incorporate your knowledge and uncertainty in the underlying model. But how to do this?

## References

This dataset and some of this code is taken from a nice online book, [Probabilistic Programming and Bayesian Methods for Hackers](#).

## Next time:

Sampling from the posterior!